

# C4\_W2\_lecture\_notebook\_Transformer\_Decoder

October 29, 2020

## 1 The Transformer Decoder: Ungraded Lab Notebook

In this notebook, you'll explore the transformer decoder and how to implement it with Trax.

### 1.1 Background

In the last lecture notebook, you saw how to translate the mathematics of attention into NumPy code. Here, you'll see how multi-head causal attention fits into a GPT-2 transformer decoder, and how to build one with Trax layers. In the assignment notebook, you'll implement causal attention from scratch, but here, you'll exploit the handy-dandy `tl.CausalAttention()` layer.

The schematic below illustrates the components and flow of a transformer decoder. Note that while the algorithm diagram flows from the bottom to the top, the overview and subsequent Trax layer codes are top-down.

### 1.2 Imports

```
[1]: import sys
import os

import time
import numpy as np
import gin

import textwrap
wrapper = textwrap.TextWrapper(width=70)

import trax
from trax import layers as tl
from trax.fastmath import numpy as jnp

# to print the entire np array
np.set_printoptions(threshold=sys.maxsize)
```

```
INFO:tensorflow:tokens_length=568 inputs_length=512 targets_length=114
noise_density=0.15 mean_noise_span_length=3.0
```

### 1.3 Sentence gets embedded, add positional encoding

Embed the words, then create vectors representing each word's position in each sentence  $\in \{0, 1, 2, \dots, K\} = \text{range}(\text{max\_len})$ , where  $\text{max\_len} = K + 1$ )

```
[2]: def PositionalEncoder(vocab_size, d_model, dropout, max_len, mode):
    """Returns a list of layers that:
    1. takes a block of text as input,
    2. embeds the words in that text, and
    3. adds positional encoding,
       i.e. associates a number in range(max_len) with
       each word in each sentence of embedded input text

    The input is a list of tokenized blocks of text

    Args:
        vocab_size (int): vocab size.
        d_model (int): depth of embedding.
        dropout (float): dropout rate (how much to drop out).
        max_len (int): maximum symbol length for positional encoding.
        mode (str): 'train' or 'eval'.
    """
    # Embedding inputs and positional encoder
    return [
        # Add embedding layer of dimension (vocab_size, d_model)
        tl.Embedding(vocab_size, d_model),
        # Use dropout with rate and mode specified
        tl.Dropout(rate=dropout, mode=mode),
        # Add positional encoding layer with maximum input length and mode
        ↪specified
        tl.PositionalEncoding(max_len=max_len, mode=mode)]
```

### 1.4 Multi-head causal attention

The layers and array dimensions involved in multi-head causal attention (which looks at previous words in the input text) are summarized in the figure below:

`tl.CausalAttention()` does all of this for you! You might be wondering, though, whether you need to pass in your input text 3 times, since for causal attention, the queries Q, keys K, and values V all come from the same source. Fortunately, `tl.CausalAttention()` handles this as well by making use of the `tl.Branch()` combinator layer. In general, each branch within a `tl.Branch()` layer performs parallel operations on copies of the layer's inputs. For causal attention, each branch (representing Q, K, and V) applies a linear transformation (i.e. a dense layer without a subsequent activation) to its copy of the input, then splits that result into heads. You can see the syntax for this in the screenshot from the `trax.layers.attention.py` [source code](#) below:

## 1.5 Feed-forward layer

- Typically ends with a ReLU activation, but we'll leave open the possibility of a different activation
- Most of the parameters are here

```
[3]: def FeedForward(d_model, d_ff, dropout, mode, ff_activation):  
    """Returns a list of layers that implements a feed-forward block.  
  
    The input is an activation tensor.  
  
    Args:  
        d_model (int): depth of embedding.  
        d_ff (int): depth of feed-forward layer.  
        dropout (float): dropout rate (how much to drop out).  
        mode (str): 'train' or 'eval'.  
        ff_activation (function): the non-linearity in feed-forward layer.  
  
    Returns:  
        list: list of trax.layers.combinators.Serial that maps an activation_  
→tensor to an activation tensor.  
    """  
  
    # Create feed-forward block (list) with two dense layers with dropout and_  
→input normalized  
    return [  
        # Normalize layer inputs  
        tl.LayerNorm(),  
        # Add first feed forward (dense) layer (don't forget to set the correct_  
→value for n_units)  
        tl.Dense(d_ff),  
        # Add activation function passed in as a parameter (you need to call it!  
→)  
        ff_activation(), # Generally ReLU  
        # Add dropout with rate and mode specified (i.e., don't use dropout_  
→during evaluation)  
        tl.Dropout(rate=dropout, mode=mode),  
        # Add second feed forward layer (don't forget to set the correct value_  
→for n_units)  
        tl.Dense(d_model),  
        # Add dropout with rate and mode specified (i.e., don't use dropout_  
→during evaluation)  
        tl.Dropout(rate=dropout, mode=mode)  
    ]
```

## 1.6 Decoder block

Here, we return a list containing two residual blocks. The first wraps around the causal attention layer, whose inputs are normalized and to which we apply dropout regulation. The second wraps around the feed-forward layer. You may notice that the second call to `tl.Residual()` doesn't call a normalization layer before calling the feed-forward layer. This is because the normalization layer is included in the feed-forward layer.

```
[4]: def DecoderBlock(d_model, d_ff, n_heads,
                    dropout, mode, ff_activation):
    """Returns a list of layers that implements a Transformer decoder block.

    The input is an activation tensor.

    Args:
        d_model (int): depth of embedding.
        d_ff (int): depth of feed-forward layer.
        n_heads (int): number of attention heads.
        dropout (float): dropout rate (how much to drop out).
        mode (str): 'train' or 'eval'.
        ff_activation (function): the non-linearity in feed-forward layer.

    Returns:
        list: list of trax.layers.combinators.Serial that maps an activation_
        ↪ tensor to an activation tensor.
    """

    # Add list of two Residual blocks: the attention with normalization and_
    ↪ dropout and feed-forward blocks
    return [
        tl.Residual(
            # Normalize layer input
            tl.LayerNorm(),
            # Add causal attention
            tl.CausalAttention(d_model, n_heads=n_heads, dropout=dropout, ↪
            ↪ mode=mode)
        ),
        tl.Residual(
            # Add feed-forward block
            # We don't need to normalize the layer inputs here. The feed-forward_
            ↪ block takes care of that for us.
            FeedForward(d_model, d_ff, dropout, mode, ff_activation)
        ),
    ]
```

## 1.7 The transformer decoder: putting it all together

### 1.8 A.k.a. repeat N times, dense layer and softmax for output

```
[5]: def TransformerLM(vocab_size=33300,
                        d_model=512,
                        d_ff=2048,
                        n_layers=6,
                        n_heads=8,
                        dropout=0.1,
                        max_len=4096,
                        mode='train',
                        ff_activation=tl.Relu):
    """Returns a Transformer language model.

    The input to the model is a tensor of tokens. (This model uses only the
    decoder part of the overall Transformer.)

    Args:
        vocab_size (int): vocab size.
        d_model (int): depth of embedding.
        d_ff (int): depth of feed-forward layer.
        n_layers (int): number of decoder layers.
        n_heads (int): number of attention heads.
        dropout (float): dropout rate (how much to drop out).
        max_len (int): maximum symbol length for positional encoding.
        mode (str): 'train', 'eval' or 'predict', predict mode is for fast_
    → inference.
        ff_activation (function): the non-linearity in feed-forward layer.

    Returns:
        trax.layers.combinators.Serial: A Transformer language model as a layer_
    → that maps from a tensor of tokens
        to activations over a vocab set.
    """

    # Create stack (list) of decoder blocks with n_layers with necessary_
    → parameters
    decoder_blocks = [
        DecoderBlock(d_model, d_ff, n_heads, dropout, mode, ff_activation) for_
    → _ in range(n_layers)]

    # Create the complete model as written in the figure
    return tl.Serial(
        # Use teacher forcing (feed output of previous step to current step)
        tl.ShiftRight(mode=mode),
        # Add embedding inputs and positional encoder
```

```

        PositionalEncoder(vocab_size, d_model, dropout, max_len, mode),
        # Add decoder blocks
        decoder_blocks,
        # Normalize layer
        tl.LayerNorm(),

        # Add dense layer of vocab_size (since need to select a word to
→translate to)
        # (a.k.a., logits layer. Note: activation already set by ff_activation)
        tl.Dense(vocab_size),
        # Get probabilities with Logsoftmax
        tl.LogSoftmax()
    )

```

## 1.9 Concluding remarks

In this week's assignment, you'll see how to train a transformer decoder on the [cnn\\_dailymail](#) dataset, available from TensorFlow Datasets (part of TensorFlow Data Services). Because training such a model from scratch is time-intensive, you'll use a pre-trained model to summarize documents later in the assignment. Due to time and storage concerns, we will also not train the decoder on a different summarization dataset in this lab. If you have the time and space, we encourage you to explore the other [summarization](#) datasets at TensorFlow Datasets. Which of them might suit your purposes better than the `cnn_dailymail` dataset? Where else can you find datasets for text summarization models?